

NATIONAL INSTITUTE OF TECHNOLOGY KARNATAKA SURATHKAL
DEPARTMENT OF INFORMATION TECHNOLOGY
IT301 : Parallel Computing Lab
PC Lab 9 [Total marks = 10]
Date 26 October 2021

CUDA Programs in Google Colab

Goto -> <https://colab.research.google.com/notebooks/intro.ipynb>

Open New Notebook.

Change Run time type to GPU

Setup the environment for running CUDA program as given in following link.

<https://www.geeksforgeeks.org/how-to-run-cuda-c-c-on-jupyter-notebook-in-google-colaboratory/>

After setting up the notebook for running CUDA execute the following programs.

Program 1: To know details of the device. Run the program and Explain the result.
[2 marks]

```
%%CU
#include<stdio.h>
int main()
{
    int devcount;
    cudaGetDeviceCount(&devcount);
    printf("Device count:%d\n",devcount);
    for (int i = 0; i < devcount; ++i)
    {
        // Get device properties
        printf("\nCUDA Device # %d\n", i);
        cudaDeviceProp devProp;
        cudaGetDeviceProperties(&devProp, i);
        printf("Name:%s\n", devProp.name);
        printf("Compute capability: %d.%d\n",devProp.major,devProp.minor);
        printf("Warp Size %d\n",devProp.warpSize);
        printf("Total global memory:%u bytes\n",devProp.totalGlobalMem);
        printf("Total shared memory per block: %u bytes\n", devProp.sharedMemPerBlock);
        printf("Total registers per block : %d\n",devProp.regsPerBlock);
        printf("Clock rate: %d khz\n",devProp.clockRate);
        printf("Maximum threads per block:%d\n", devProp.maxThreadsPerBlock);
        for (int i = 0; i < 3; ++i)
        printf("Maximum dimension %d of block: %d\n", i, devProp.maxThreadsDim[i]);
        for (int i = 0; i <= 2; ++i)
        printf("Maximum dimension %d of grid: %d\n", i, devProp.maxGridSize[i]);
        printf("Number of multiprocessors:%d\n", devProp.multiProcessorCount);
    }
    return 0;
}
```

Program 2: Hello world program. Record the result and write the observation.
[2 marks]

```
%%CU
#include<stdio.h>
#include<cuda.h>
__global__ void helloworld(void)
{
printf("Hello World from GPU\n");
}
int main() {
helloworld<<<1,10>>>();
printf("Hello World\n");
return 0;
}
```

Program 3: Program to perform $c[i] = a[i] + b[i]$; Here, $c[i]$ is calculated for all i . But results are displayed only for few $c[i]$. Explain your observation. [2 x 3 = 6 Marks]
Run the program for following and note down the time.

- a) `vecAdd<<<1,100>>>(d_a, d_b, d_c, n);`
- b) `vecAdd<<<1,50>>>(d_a, d_b, d_c, n);`
- c) `vecAdd<<<2,50>>>(d_a, d_b, d_c, n);`

```
%%CU
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <sys/time.h>

__global__ void vecAdd(double *a, double *b, double *c, int n)
{
// Get global thread
int id = blockIdx.x*blockDim.x+threadIdx.x;
// Do not go out of bounds
if (id < n)
c[id] = a[id] + b[id];
}
int main( int argc, char* argv[] )
{
// Size of vectors
int n = 100;
//time variables
struct timeval t1, t2;
// Host input vectors
double *h_a, *h_b;
//Host output vector
double *h_c;
// Device input vectors
```

```

double *d_a, *d_b;
//Device output vector
double *d_c;
// Size, in bytes, of each vector
size_t bytes = n*sizeof(double);
// Allocate memory for each vector on host
h_a = (double*)malloc(bytes);
h_b = (double*)malloc(bytes);
h_c = (double*)malloc(bytes);
// Allocate memory for each vector on GPU
cudaMalloc(&d_a, bytes);
cudaMalloc(&d_b, bytes);
cudaMalloc(&d_c, bytes);
int i;
// Initialize vectors on host
for( i = 0; i < n; i++ ) {
h_a[i] = i+1;
h_b[i] = i+1;
}
// Copy host vectors to device
cudaMemcpy( d_a, h_a, bytes, cudaMemcpyHostToDevice);
cudaMemcpy( d_b, h_b, bytes, cudaMemcpyHostToDevice);

gettimeofday(&t1, 0);

// Execute the kernel
vecAdd<<<1,100>>>>(d_a, d_b, d_c, n);
cudaDeviceSynchronize();
gettimeofday(&t2, 0);

// Copy array back to host
cudaMemcpy( h_c, d_c, bytes, cudaMemcpyDeviceToHost );

for(i=0; i<n; i=i+10)
printf("c[%d]=%f\n",i,h_c[i]);
double time = (1000000.0*(t2.tv_sec-t1.tv_sec) + t2.tv_usec-t1.tv_usec)/1000.0;

printf("Time to generate: %3.10f ms \n", time);

// Release device memory
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);
// Release host memory
free(h_a);
free(h_b);
free(h_c);
return 0;
}

```